

```
In [1]: import pandas as pd
import numpy as np
import os
import matplotlib.pyplot as plt
from matplotlib import pyplot
import seaborn as sns
from sklearn.decomposition import PCA
from sklearn.preprocessing import scale
from statsmodels.tsa.arima_model import ARIMA
from datetime import datetime
```

```
In [2]: #insert data from csv
df = pd.read_csv('/Users/amrita/Desktop/zillow data.csv')
```

```
In [41]: df.head()
```

Out[41]:

	RegionID	SizeRank	RegionName	RegionType	StateName	1/31/00	2/29/00	3/31/00	4/30/00
0	102001	0	United States	Country	NaN	127104.0	127448.0	127809.0	128544.0
1	394913	1	New York, NY	Msa	NY	223875.0	225213.0	226416.0	228781.0
2	753899	2	Los Angeles-Long Beach-Anaheim, CA	Msa	CA	231151.0	231956.0	233189.0	235531.0
3	394463	3	Chicago, IL	Msa	IL	169017.0	169416.0	169932.0	170961.0
4	394514	4	Dallas-Fort Worth, TX	Msa	TX	130276.0	130380.0	130466.0	130671.0

5 rows × 268 columns

```
In [42]: #transpose data
ndf= df.melt(id_vars=["RegionID", "SizeRank", "RegionName", "RegionType", "StateName"],
             var_name="Date",
             value_name="Price")
```

```
In [43]: ndf.head()
```

Out[43]:

	RegionID	SizeRank	RegionName	RegionType	StateName	Date	Price
0	102001	0	United States	Country	NaN	1/31/00	127104.0
1	394913	1	New York, NY	Msa	NY	1/31/00	223875.0
2	753899	2	Los Angeles-Long Beach-Anaheim, CA	Msa	CA	1/31/00	231151.0
3	394463	3	Chicago, IL	Msa	IL	1/31/00	169017.0
4	394514	4	Dallas-Fort Worth, TX	Msa	TX	1/31/00	130276.0

```
In [44]: #summary of the datatype
ndf.dtypes
```

```
Out[44]: RegionID      int64
SizeRank      int64
RegionName     object
RegionType     object
StateName      object
Date           object
Price         float64
dtype: object
```

```
In [45]: #no. of missing values by column
ndf.isna().sum()
```

```
Out[45]: RegionID      0
SizeRank      0
RegionName     0
RegionType     0
StateName     263
Date           0
Price         48727
dtype: int64
```

```
In [8]: #Checking for the total count of Region name=United States
```

```
i= ndf[ndf['RegionName']=='United States']
i.shape
```

```
Out[8]: (263, 7)
```

```
In [46]: i.head()
```

```
Out[46]:
```

	RegionID	SizeRank	RegionName	RegionType	StateName	Date	Price
0	102001	0	United States	Country	NaN	1/31/00	127104.0
908	102001	0	United States	Country	NaN	2/29/00	127448.0
1816	102001	0	United States	Country	NaN	3/31/00	127809.0
2724	102001	0	United States	Country	NaN	4/30/00	128546.0
3632	102001	0	United States	Country	NaN	5/31/00	129288.0

In [47]: *#Dropping all rows with Region Name=United States and creating a new dataframe*

```
df_new= ndf[ndf.RegionName != 'United States']
df_new.head()
```

Out[47]:

	RegionID	SizeRank	RegionName	RegionType	StateName	Date	Price
1	394913	1	New York, NY	Msa	NY	1/31/00	223875.0
2	753899	2	Los Angeles-Long Beach-Anaheim, CA	Msa	CA	1/31/00	231151.0
3	394463	3	Chicago, IL	Msa	IL	1/31/00	169017.0
4	394514	4	Dallas-Fort Worth, TX	Msa	TX	1/31/00	130276.0
5	394974	5	Philadelphia, PA	Msa	PA	1/31/00	129615.0

In [69]: df_new.describe()

Out[69]:

	RegionID	SizeRank	Price
count	238541.000000	238541.000000	1.898140e+05
mean	415361.502756	458.604190	1.612216e+05
std	83488.890005	267.525459	1.038787e+05
min	394297.000000	1.000000	2.848100e+04
25%	394548.000000	227.000000	9.925225e+04
50%	394804.000000	455.000000	1.320360e+05
75%	395050.000000	687.000000	1.877395e+05
max	753929.000000	933.000000	1.506129e+06

In [78]: df_new.(RegionName=='New York, NY').Date.unique()

```
File "<ipython-input-78-9002fa9f8b66>", line 1
    df_new.(RegionName=='New York, NY').Date.unique()
           ^
SyntaxError: invalid syntax
```

In [11]: df_new.shape

Out[11]: (238541, 7)

```
In [12]: df_new.dtypes
```

```
Out[12]: RegionID      int64
SizeRank      int64
RegionName     object
RegionType     object
StateName      object
Date           object
Price         float64
dtype: object
```

```
In [48]: #Changing date from object to datetime
```

```
df_new['Date'] = pd.to_datetime(df_new['Date'])
df_new.head()
```

/Users/amrita/opt/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

This is separate from the ipykernel package so we can avoid doing imports until

```
Out[48]:
```

	RegionID	SizeRank	RegionName	RegionType	StateName	Date	Price
1	394913	1	New York, NY	Msa	NY	2000-01-31	223875.0
2	753899	2	Los Angeles-Long Beach-Anaheim, CA	Msa	CA	2000-01-31	231151.0
3	394463	3	Chicago, IL	Msa	IL	2000-01-31	169017.0
4	394514	4	Dallas-Fort Worth, TX	Msa	TX	2000-01-31	130276.0
5	394974	5	Philadelphia, PA	Msa	PA	2000-01-31	129615.0

```
In [49]: df_new.dtypes
```

```
Out[49]: RegionID      int64
SizeRank      int64
RegionName     object
RegionType     object
StateName      object
Date          datetime64[ns]
Price         float64
dtype: object
```

In [50]:

```
#no. of missing values by column in the new dataset

df_new.isna().sum()
```

```
Out[50]: RegionID          0
         SizeRank        0
         RegionName       0
         RegionType       0
         StateName        0
         Date            0
         Price          48727
         dtype: int64
```

In [51]: *#percent missing values by each column*

```
percent_missing = df_new.isnull().sum() * 100 / len(df_new)
missing_value_df = pd.DataFrame({'column_name': df_new.columns,
                                'percent_missing': percent_missing})
missing_value_df.sort_values('percent_missing', inplace=True)
print(missing_value_df)
```

	column_name	percent_missing
RegionID	RegionID	0.000000
SizeRank	SizeRank	0.000000
RegionName	RegionName	0.000000
RegionType	RegionType	0.000000
StateName	StateName	0.000000
Date	Date	0.000000
Price	Price	20.427096

In [52]: *#impute missing values using interpolation method*

```
interpolated = df_new.interpolate(method='linear')
interpolated.head()
```

Out[52]:

	RegionID	SizeRank	RegionName	RegionType	StateName	Date	Price
1	394913	1	New York, NY	Msa	NY	2000-01-31	223875.0
2	753899	2	Los Angeles-Long Beach-Anaheim, CA	Msa	CA	2000-01-31	231151.0
3	394463	3	Chicago, IL	Msa	IL	2000-01-31	169017.0
4	394514	4	Dallas-Fort Worth, TX	Msa	TX	2000-01-31	130276.0
5	394974	5	Philadelphia, PA	Msa	PA	2000-01-31	129615.0

In [54]: *#no missing values*

```
interpolated.isna().sum()
```

Out[54]:

RegionID	0
SizeRank	0
RegionName	0
RegionType	0
StateName	0
Date	0
Price	0

dtype: int64

In [55]: *#percent missing values by each column*

```
percentage_missing = interpolated.isnull().sum() * 100 / len(interpolated)
missing_value = pd.DataFrame({'column_name': interpolated.columns,
                              'percent_missing': percentage_missing})
missing_value.sort_values('percent_missing', inplace=True)
print(missing_value)
```

	column_name	percent_missing
RegionID	RegionID	0.0
SizeRank	SizeRank	0.0
RegionName	RegionName	0.0
RegionType	RegionType	0.0
StateName	StateName	0.0
Date	Date	0.0
Price	Price	0.0

In [56]: *# total unique region names*

```
interpolated.RegionName.nunique()
```

Out[56]: 907

In [57]: interpolated['RegionName'].value_counts().head(100)

Out[57]:

Albemarle, NC	263
Hereford, TX	263
Scottsbluff, NE	263
Red Bluff, CA	263
Indianapolis, IN	263
...	
Fort Morgan, CO	263
Worcester, MA	263
Menomonie, WI	263
Ocean City, NJ	263
Eau Claire, WI	263

Name: RegionName, Length: 100, dtype: int64

```
In [66]: #Distribution of price by RegionName  
region_price_means = interpolated.groupby("RegionName")["Price"].mean()  
region_price_means.head(90)
```

Out[66]:

Price	
RegionName	
Aberdeen, SD	131921.124585
Aberdeen, WA	194096.368061
Abilene, TX	134786.818931
Ada, OK	80116.889734
Adrian, MI	133173.357414
...	...
Big Stone Gap, VA	90701.484791
Billings, MT	189141.785171
Binghamton, NY	106751.771863
Birmingham, AL	147310.996198
Bismarck, ND	184398.741445

90 rows × 1 columns

Average Price by Region Name

Price

RegionName

Price (\$)



```
In [59]: # statistical summary of all the numerical columns
```

```
interpolated.describe().T
```

Out[59]:

	count	mean	std	min	25%	50%	75%	ma
RegionID	238541.0	415361.502756	83488.890005	394297.0	394548.0	394804.0	395050.0	753929.
SizeRank	238541.0	458.604190	267.525459	1.0	227.0	455.0	687.0	933.
Price	238541.0	154240.877113	98655.340633	28481.0	95525.0	126544.0	179056.0	1506129.

```
In [60]: interpolated.Price.hist(bins=20)
plt.xlabel('Price')
plt.ylabel('Count')
plt.title('Distribution of Price');
```



```
In [61]: interpolated['log_base10'] = np.log10(interpolated['Price'])
interpolated.head()
```

Out[61]:


	RegionID	SizeRank	RegionName	RegionType	StateName	Date	Price	log_base10
1	394913	1	New York, NY	Msa	NY	2000-01-31	223875.0	5.350006
2	753899	2	Los Angeles-Long Beach-Anaheim, CA	Msa	CA	2000-01-31	231151.0	5.363896
3	394463	3	Chicago, IL	Msa	IL	2000-01-31	169017.0	5.227930
4	394514	4	Dallas-Fort Worth, TX	Msa	TX	2000-01-31	130276.0	5.114864
5	394974	5	Philadelphia, PA	Msa	PA	2000-01-31	129615.0	5.112655

```
In [62]: #taking natural log to achieve a normal distribution
interpolated['natural_log'] = np.log(interpolated['Price'])
interpolated.head()
```

Out[62]:

	RegionID	SizeRank	RegionName	RegionType	StateName	Date	Price	log_base10	natural
1	394913	1	New York, NY	Msa	NY	2000-01-31	223875.0	5.350006	12.318
2	753899	2	Los Angeles-Long Beach-Anaheim, CA	Msa	CA	2000-01-31	231151.0	5.363896	12.350
3	394463	3	Chicago, IL	Msa	IL	2000-01-31	169017.0	5.227930	12.037
4	394514	4	Dallas-Fort Worth, TX	Msa	TX	2000-01-31	130276.0	5.114864	11.777
5	394974	5	Philadelphia, PA	Msa	PA	2000-01-31	129615.0	5.112655	11.772

```
In [63]: interpolated.natural_log.hist(bins=20)
plt.xlabel('Price')
plt.ylabel('Count')
plt.title('Distribution of Price');
```



```
In [64]: interpolated.drop('log_base10', axis=1, inplace=True)
```

```
In [65]: interpolated.head()
```

```
Out[65]:
```

	RegionID	SizeRank	RegionName	RegionType	StateName	Date	Price	natural_log
1	394913	1	New York, NY	Msa	NY	2000-01-31	223875.0	12.318843
2	753899	2	Los Angeles-Long Beach-Anaheim, CA	Msa	CA	2000-01-31	231151.0	12.350826
3	394463	3	Chicago, IL	Msa	IL	2000-01-31	169017.0	12.037755
4	394514	4	Dallas-Fort Worth, TX	Msa	TX	2000-01-31	130276.0	11.777411
5	394974	5	Philadelphia, PA	Msa	PA	2000-01-31	129615.0	11.772324

```
In [67]: interpolated.shape
```

```
Out[67]: (238541, 8)
```

```
In [68]: interpolated.to_excel('output1.xlsx', engine='xlsxwriter')
```